

# CS4.301 Data & Applications

Ponnurangam Kumaraguru ("PK")  
#ProfGiri @ IIIT Hyderabad



pk.profgiri



/in/ponguru



@ponguru



Ponnurangam.kumaraguru

## Nested Queries (cont'd.)

```
SELECT  DISTINCT Pnumber
FROM    PROJECT
WHERE   Pnumber IN
        (SELECT  Pnumber
         FROM    PROJECT, DEPARTMENT, EMPLOYEE
         WHERE   Dnum = Dnumber AND
                 Mgr_ssn = Ssn and Lname = 'Smith')
OR
        Pnumber IN
        (SELECT  Pno
         FROM    WORKS_ON, EMPLOYEE
         WHERE   Essn = Ssn AND Lname = 'Smith');
```

Pnumber
1
2

2 rows in set (0.01 sec)

# Nested Queries (cont'd.)

Use tuples of values in comparisons

Place them within parentheses

Select distinct essn

From works\_on

Where (pno, hours) IN

(Select pno, hours from  
works\_on where essn =  
'123456789');

```
mysql> Select pno, hours from works_on where essn = '123456789';
+-----+-----+
| pno | hours |
+-----+-----+
| 1   | 32.5  |
| 2   | 7.5   |
+-----+-----+
2 rows in set (0.04 sec)
```

```
mysql> Select distinct essn
-> From works_on
-> Where (pno, hours) IN (Select pno, hours from works_on where essn = '123456789');
+-----+
| essn |
+-----+
| 123456789 |
+-----+
1 row in set (0.01 sec)
```

# Nested Queries (cont'd.)

Use other comparison operators to compare a single value  $v$

= ANY (or = SOME) operator

Returns TRUE if the value  $v$  is equal to some value in the set  $V$  and is hence equivalent to IN

Other operators that can be combined with ANY (or SOME): >, >=, <, <=, and <>

ALL: value must exceed all values from nested query

Select lname, fname,  
salary from employee  
where salary > all (select  
salary from employee  
where dno = 5);

```
mysql> Select lname, fname, salary from employee w
here salary > all (select salary from employee whe
re dno = 5);
+-----+-----+-----+
| lname | fname | salary |
+-----+-----+-----+
| Borg  | James | 55000  |
| Wallace | Jennifer | 43000  |
+-----+-----+-----+
2 rows in set (0.00 sec)
```

```
mysql> select salary from employee where dno = 5;
+-----+
| salary |
+-----+
| 30000  |
| 40000  |
| 25000  |
| 38000  |
+-----+
4 rows in set (0.00 sec)
```

## Nested Queries (cont'd.)

Avoid potential errors and ambiguities

Create tuple variables (aliases) for all tables referenced in SQL query

**Query 16.** Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

Select e.fname, e.lname from  
employee as e where e.ssn in  
(select essn from dependent as  
d where  
e.fname=d.dependent\_name  
and e.sex=d.sex);

```
mysql> Select e.fname, e.lname from employee as e w  
here e.ssn in (select essn from dependent as d wher  
e e.fname=d.dependent_name and e.sex=d.sex);  
Empty set (0.00 sec)
```

# Correlated Nested Queries

**Queries that are nested using the = or IN comparison operator** can be collapsed into one single block, last query can be changed like ??? Ideas?

```
Select e.fname, e.lname from  
employee as e where e.ssn in  
(select essn from dependent as  
d where  
e.fname=d.dependent_name  
and e.sex=d.sex);
```

**Correlated** nested query

Evaluated once for each tuple in the outer query

# Correlated Nested Queries

**Queries that are nested using the = or IN comparison operator** can be collapsed into one single block, last query can be changed like

```
Select e.fname, e.lname from  
employee as e where e.ssn in  
(select essn from dependent as  
d where  
e.fname=d.dependent_name  
and e.sex=d.sex);
```

```
SELECT E.Fname, E.Lname  
FROM EMPLOYEE AS E,  
DEPENDENT AS D WHERE  
E.Ssn=D.Essn AND E.Sex=D.Sex  
AND  
E.Fname=D.Dependent_name;
```

**Correlated** nested query

Evaluated once for each tuple in the outer query

# Explicit Sets and Renaming of Attributes in SQL

Can use explicit set of values in WHERE clause

```
SELECT DISTINCT Essn FROM WORKS_ON WHERE Pno IN (1, 2, 3);
```

```
mysql> SELECT DISTINCT Essn FROM WORKS_ON WHERE Pno|
      IN (1, 2, 3);
+-----+
| Essn   |
+-----+
| 123456789 |
| 453453453 |
| 333445555 |
| 666884444 |
+-----+
4 rows in set (0.00 sec)
```



# Explicit Sets and Renaming of Attributes in SQL

Use qualifier AS followed by desired new name

Rename any attribute that appears in the result of a query

Select e.lname as  
employee\_name, s.lname as  
supervisor\_name from employee  
as e, employee as s where  
e.super\_ssn = s.ssn;

```
mysql> Select e.lname as employee_name, s.lname as  
supervisor_name from employee as e, employee as s w  
here e.super_ssn = s.ssn;
```

employee_name	supervisor_name
Smith	Wong
Wong	Borg
English	Wong
Narayan	Wong
Wallace	Borg
Jabbar	Wallace
Zelaya	Wallace

```
7 rows in set (0.00 sec)
```

# Aggregate Functions in SQL

Used to summarize information from multiple tuples into a single-tuple summary

Built-in aggregate functions

**COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**

## **Grouping**

Create subgroups of tuples before summarizing

To select entire groups, **HAVING** clause is used

Aggregate functions can be used in the **SELECT** clause or in a **HAVING** clause

# Renaming Results of Aggregation

```
SELECT SUM(Salary),  
MAX(Salary),  
MIN(Salary), AVG(Salary)  
FROM EMPLOYEE;
```

```
mysql> SELECT SUM(Salary), MAX(Salary), MIN(Salary), AVG(Sal  
ary) FROM EMPLOYEE;  
+-----+-----+-----+-----+  
| SUM(Salary) | MAX(Salary) | MIN(Salary) | AVG(Salary) |  
+-----+-----+-----+-----+  
|      281000 |       55000 |       25000 | 35125.0000 |  
+-----+-----+-----+-----+  
1 row in set (0.00 sec)
```

# Renaming Results of Aggregation

```
SELECT SUM(Salary) AS  
Total_Sal, MAX(Salary)  
AS Highest_Sal,  
MIN(Salary) AS  
Lowest_Sal, AVG(Salary)  
AS Average_Sal FROM  
EMPLOYEE;
```

```
mysql> SELECT SUM(Salary) AS Total_Sal, MAX(Salary) AS Highest_Sal,  
MIN(Salary) AS Lowest_Sal, AVG(Salary) AS Average_Sal  
FROM EMPLOYEE;
```

Total_Sal	Highest_Sal	Lowest_Sal	Average_Sal
281000	55000	25000	35125.0000

```
1 row in set (0.00 sec)
```

## Aggregate Functions in SQL (cont'd.)

**Query 20.** Find the sum of the salaries of all employees of the 'Research' department, as well as the maximum salary, the minimum salary, and the average salary in this department.

```
SELECT SUM(Salary),  
MAX(Salary),  
MIN(Salary), AVG(Salary)  
FROM (EMPLOYEE join  
department on  
dno=dnumber) where  
dname='research';
```

```
mysql> SELECT SUM(Salary), MAX(Salary), MIN(Salary), AVG(Sal  
ary) FROM (EMPLOYEE join department on dno=dnumber) where dn  
ame='research';  
+-----+-----+-----+-----+  
| SUM(Salary) | MAX(Salary) | MIN(Salary) | AVG(Salary) |  
+-----+-----+-----+-----+  
|      133000 |        40000 |        25000 | 33250.0000 |  
+-----+-----+-----+-----+  
1 row in set (0.00 sec)
```

# Aggregate Functions in SQL (cont'd.)

**Queries 21 and 22.** Retrieve the total number of employees in the company (Q21) and the number of employees in the 'Research' department (Q22).

Select count(\*) from  
employee;

Select count(\*) from  
employee, department  
where dno=dnumber  
and dname='research';

```
mysql> Select count(*) from employee;
+-----+
| count(*) |
+-----+
|         8 |
+-----+
1 row in set (0.02 sec)

mysql> Select count(*) from employee, department where dno=dnumber and dname='research';
+-----+
| count(*) |
+-----+
|         4 |
+-----+
1 row in set (0.00 sec)
```

# Group BY example

```
SELECT Dno, COUNT(*),  
AVG(Salary) FROM  
EMPLOYEE GROUP BY  
Dno;
```

```
mysql> SELECT Dno, COUNT(*), AVG(Salary) FROM EMPLOYEE GROUP  
BY Dno;  
+-----+-----+-----+  
| Dno | COUNT(*) | AVG(Salary) |  
+-----+-----+-----+  
| 5 | 4 | 33250.0000 |  
| 1 | 1 | 55000.0000 |  
| 4 | 3 | 31000.0000 |  
+-----+-----+-----+  
3 rows in set (0.01 sec)
```

# Group BY example

```
SELECT Pnumber,  
Pname, COUNT(*) FROM  
PROJECT, WORKS_ON  
WHERE Pnumber=Pno  
GROUP BY Pname;
```

```
[mysql> SELECT Pnumber, Pname, COUNT(*) FROM PROJECT, WORKS_ON  
N WHERE Pnumber=Pno GROUP BY Pname;
```

Pnumber	Pname	COUNT(*)
10	Computerization	3
30	Newbenefits	3
1	ProductX	2
2	ProductY	3
3	ProductZ	2
20	Reorganization	3

6 rows in set (0.01 sec)



# This Lecture

# Administrativa

Guest lecture on Nov 7<sup>th</sup> evening

# Grouping: The GROUP BY and HAVING Clauses (cont'd.)

## HAVING clause

Provides a condition to select or reject an entire group:

**Query 26.** For each project *on which more than two employees work*, retrieve the project number, the project name, and the number of employees who work on the project.

```
SELECT Pnumber,  
Pname, COUNT(*) FROM  
PROJECT, WORKS_ON  
WHERE Pnumber=Pno  
GROUP BY Pnumber  
HAVING COUNT(*) > 2;
```

```
mysql> SELECT Pnumber, Pname, COUNT(*) FROM PROJECT, WORKS_ON  
N WHERE Pnumber=Pno GROUP BY Pnumber HAVING COUNT(*) > 2;  
+-----+-----+-----+  
| Pnumber | Pname           | COUNT(*) |  
+-----+-----+-----+  
|      2  | ProductY        |      3   |  
|     10  | Computerization |      3   |  
|     20  | Reorganization  |      3   |  
|     30  | Newbenefits     |      3   |  
+-----+-----+-----+  
4 rows in set (0.00 sec)
```

# EXPANDED Block Structure of SQL Queries

```
SELECT <attribute and function list>  
FROM <table list>  
[ WHERE <condition> ]  
[ GROUP BY <grouping attribute(s)> ]  
[ HAVING <group condition> ]  
[ ORDER BY <attribute list> ];
```

# Views (Virtual Tables) in SQL

## Concept of a view in SQL

Single table derived from other tables called the **defining tables**

Considered to be a virtual table that is not necessarily populated

Therefore limits update operations, no limitations in querying

In COMPANY we may frequently retrieve project name & employee name which is joining employee, works\_on, project create a view and query this single table retrieval than multiple tables

# Specification of Views in SQL

## **CREATE VIEW** command

Give table name, list of attribute names, and a query to specify the contents of the view

Create view works\_on1  
as select fname, lname,  
hours from employee,  
project, works\_on  
where ssn=essn and  
pno=pnumber;

```
mysql> Create view works_on1 as select fname, lname, hours f  
rom employee, project, works_on where ssn=essn and pno=pnumb  
er;  
Query OK, 0 rows affected (0.05 sec)
```

## Data in view, query view

```
select * from works_on1;
```

```
mysql> select * from works_on1;
+-----+-----+-----+
| fname | lname | hours |
+-----+-----+-----+
| Franklin | Wong | 10.0 |
| James | Borg | 16.0 |
| Jennifer | Wallace | 15.0 |
| Franklin | Wong | 10.0 |
| Ahmad | Jabbar | 35.0 |
| Alicia | Zelaya | 10.0 |
| Jennifer | Wallace | 20.0 |
| Ahmad | Jabbar | 5.0 |
| Alicia | Zelaya | 30.0 |
| John | Smith | 32.5 |
| Joyce | English | 20.0 |
| John | Smith | 7.5 |
| Franklin | Wong | 10.0 |
| Joyce | English | 20.0 |
| Franklin | Wong | 10.0 |
| Ramesh | Narayan | 40.0 |
+-----+-----+-----+
16 rows in set (0.01 sec)
```

```
select fname, lname from
works_on1 where
hours=10;
```

```
mysql> select fname, lname from works_on1 where hours=10;
+-----+-----+
| fname | lname |
+-----+-----+
| Franklin | Wong |
| Franklin | Wong |
| Franklin | Wong |
| Franklin | Wong |
| Alicia | Zelaya |
+-----+-----+
5 rows in set (0.01 sec)
```

## Specification of Views in SQL (cont'd.)

Once a View is defined, SQL queries can use the View relation in the FROM clause

View is always up-to-date

Responsibility of the DBMS and not the user

**DROP VIEW** command

Dispose of a view



# View Implementation, View Update, and Inline Views

Complex problem of efficiently implementing a view for querying

## **Strategy1: Query modification** approach

- Compute the view as and when needed. Do not store permanently

- Modify view query into a query on underlying base tables

- Disadvantage: inefficient for views defined via complex queries that are time-consuming to execute

# View Materialization

## **Strategy 2: View materialization**

Physically create a temporary view table when the view is first queried

Keep that table on the assumption that other queries on the view will follow

Requires efficient strategy for automatically updating the view table when the base tables are updated

## **Incremental update strategy for materialized views**

DBMS determines what new tuples must be inserted, deleted, or modified in a materialized view table

## View Materialization (contd.)

Multiple ways to handle materialization:

**immediate update** strategy updates a view as soon as the base tables are changed

**lazy update** strategy updates the view when needed by a view query

**periodic update** strategy updates the view periodically (in the latter strategy, a view query may get a result that is not up-to-date). This is commonly used in Banks, Retail store operations, etc.

# Schema Change Statements in SQL

## **Schema evolution commands**

DBA may want to change the schema while the database is operational  
Does not require recompilation of the database schema

# The DROP Command

DROP command

Used to drop named schema elements, such as tables, domains, or constraint

Drop behavior options:

CASCADE and RESTRICT

RESTRICT – schema will be dropped only if it has no elements in it

Example:

```
DROP SCHEMA COMPANY CASCADE;
```

This removes the schema and all its elements including tables, views, constraints, etc.

```
DROP TABLE DEPENDENT CASCADE;
```

If we no longer wish to track the dependents

# The DROP Command

Not only deletes all the records in the table if successful, removes the table definition from catalog

 pk.profgiri

 Ponnurangam.kumaraguru

 /in/ponguru

 ponguru

 pk.guru@iiit.ac.in

Thank you  
for attending  
the class!!!